



# **Machine Learning 410**

## **Lesson 12**

### **Introduction to Neural Machine Translation with RNNs**

Steve Elston

# Reminders


- Discussion – get an easy 5 points!
  - Discussion 8 closes Tonight, Nov 21.
  - Discussion 9 closes Thursday, Dec 5.
- Homework – **Updated homework is in Canvas**
  - Homework 7 due December 2

# Outline

- Review of recurrent neural network (RNN) architectures
- The long-short term memory (LSTM) unit
- Bidirectional RNNs (BiRNN)
- Introduction to statistical machine translation (SMT)
- Evaluation of STM models
- Overview of neural machine translation (NMT)
- The attention mechanism for NMT
- Google translate


# RNN Architectures

## Feedforward network




# RNN Architectures

Generative model where input vector generates output sequence




# RNN Architectures

Classification of input sequence




# RNN Architectures

Sequence-to-sequence model – Generative



# RNN Architectures

Sequence-to-sequence model with **context** and embedding






# Sequence Generation with RNNs

- Given an input value  $\mathbf{x}$ , we want to generate an output sequence  $\mathbf{y}$ .
  - Activation is function of input and last output:  $Y_t = h_{t-1}(y_{t-1}, x)$
  - Is a generative model
- Train by minimizing the loss with respect to the desired response
  - Loss function =  $J(U, V, W)$
- Applications:
  - Response to question – chat bots
  - Caption images
  - Machine translation

# Sequence Generation with RNNs



# Long-Short Term Memory

- Problems training simple recurrent architectures have led to the development of better approaches
  - Vanishing and exploding gradients common
  - Consider the following recurrent relationship:

$$s^{(t+n)} = f(f(\dots f(f(s^{(t)}; \theta)); \theta))$$

For a long recurrence, large n:

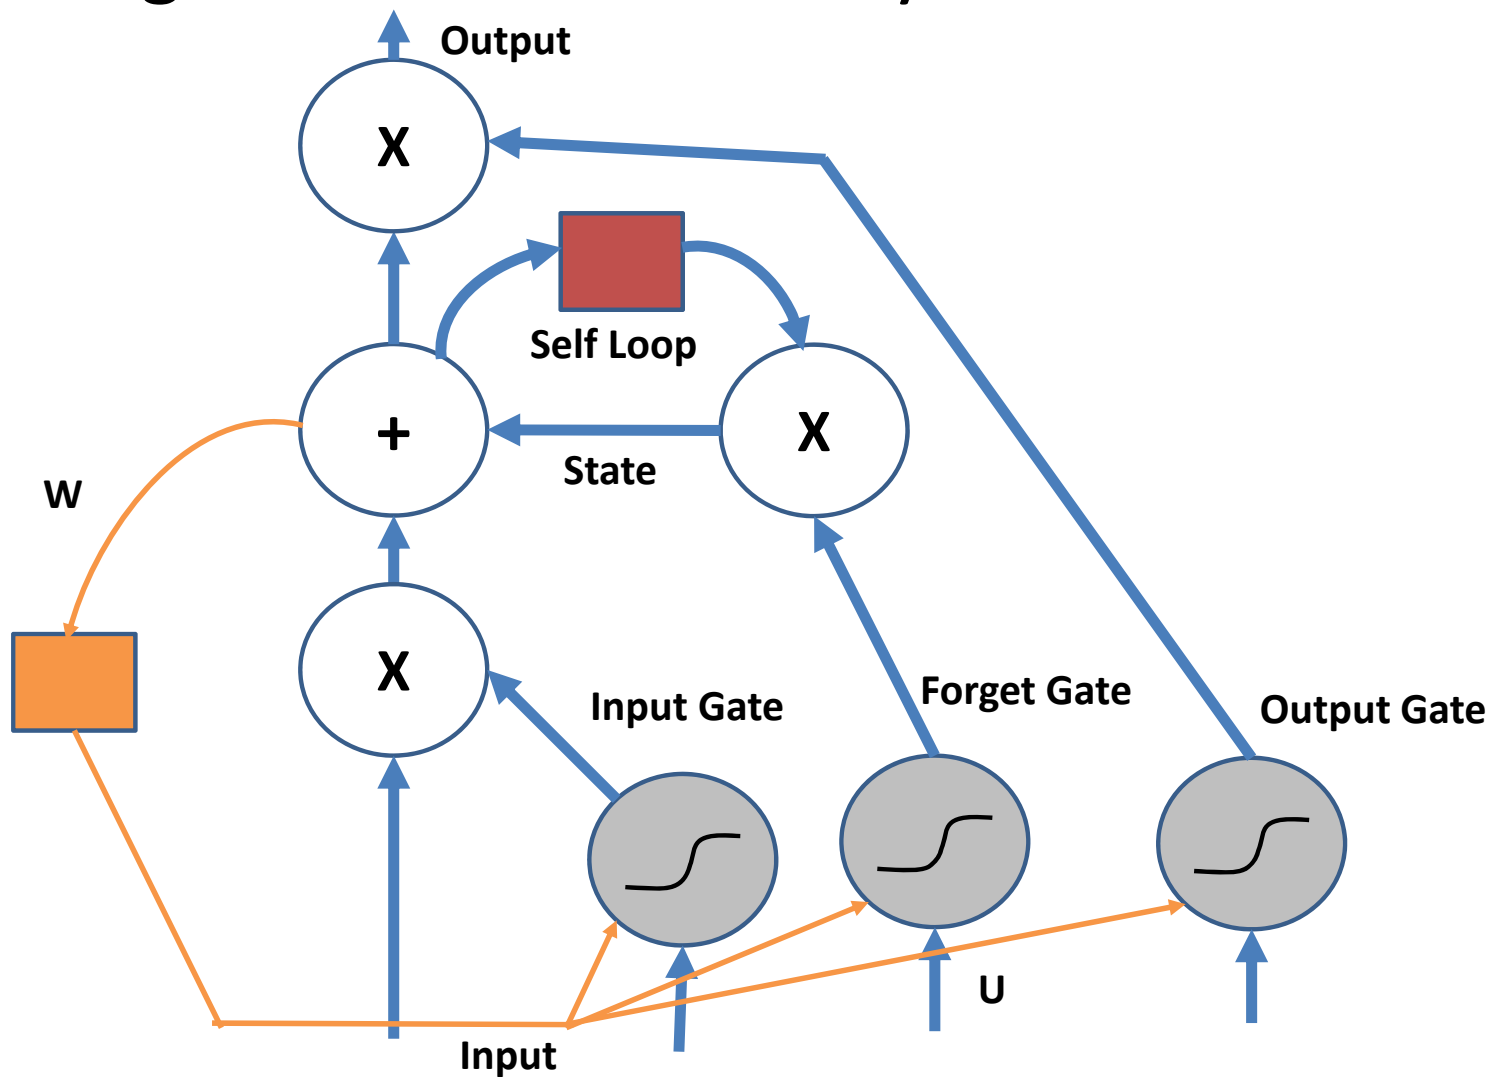
If  $f(s; \Theta) > 1.0$  the gradient grows exponentially

If  $f(s; \Theta) < 1.0$  the gradient vanishes

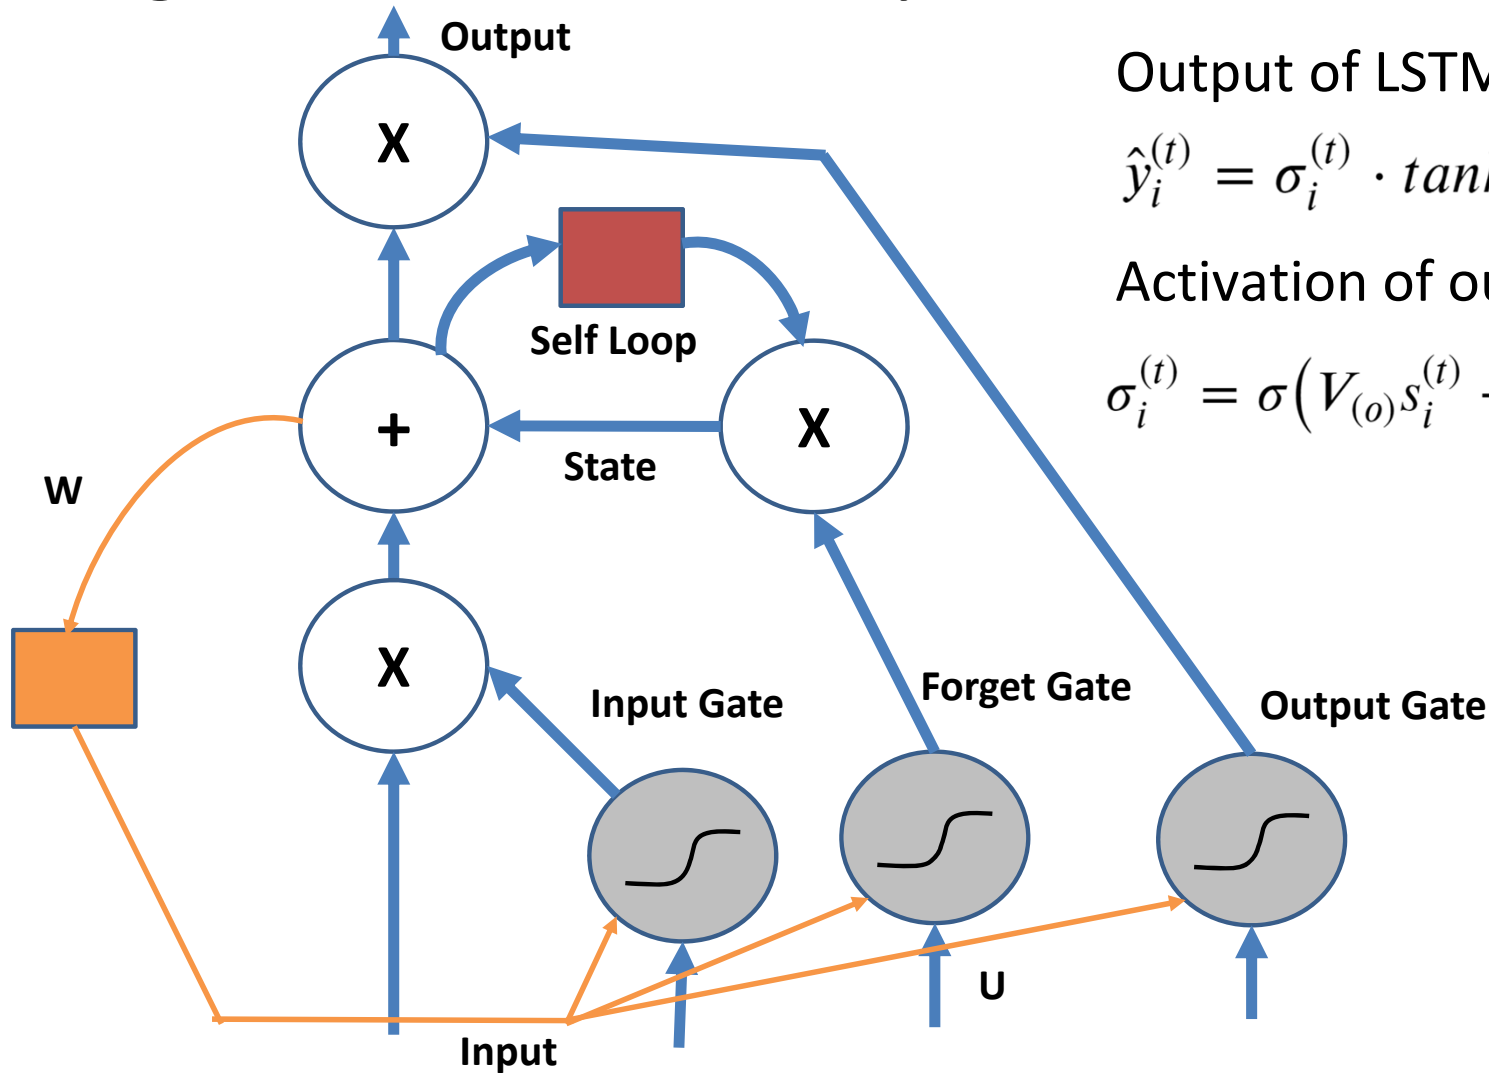
# Long-Short Term Memory

- How can one create a neural net for modeling sequences with stable gradients?
  - **Memory** lets the NN operate at multiple time scales
  - **Forget gates** break the recurrence relationship and stabilize the gradient
- The **Long-Short Term Memory (LSTM)** neural network was an early architecture using memory and forget gates
- LSTM used for speech recognition, handwriting generation, machine translation, etc.

# Long-Short Term Memory



# Long-Short Term Memory



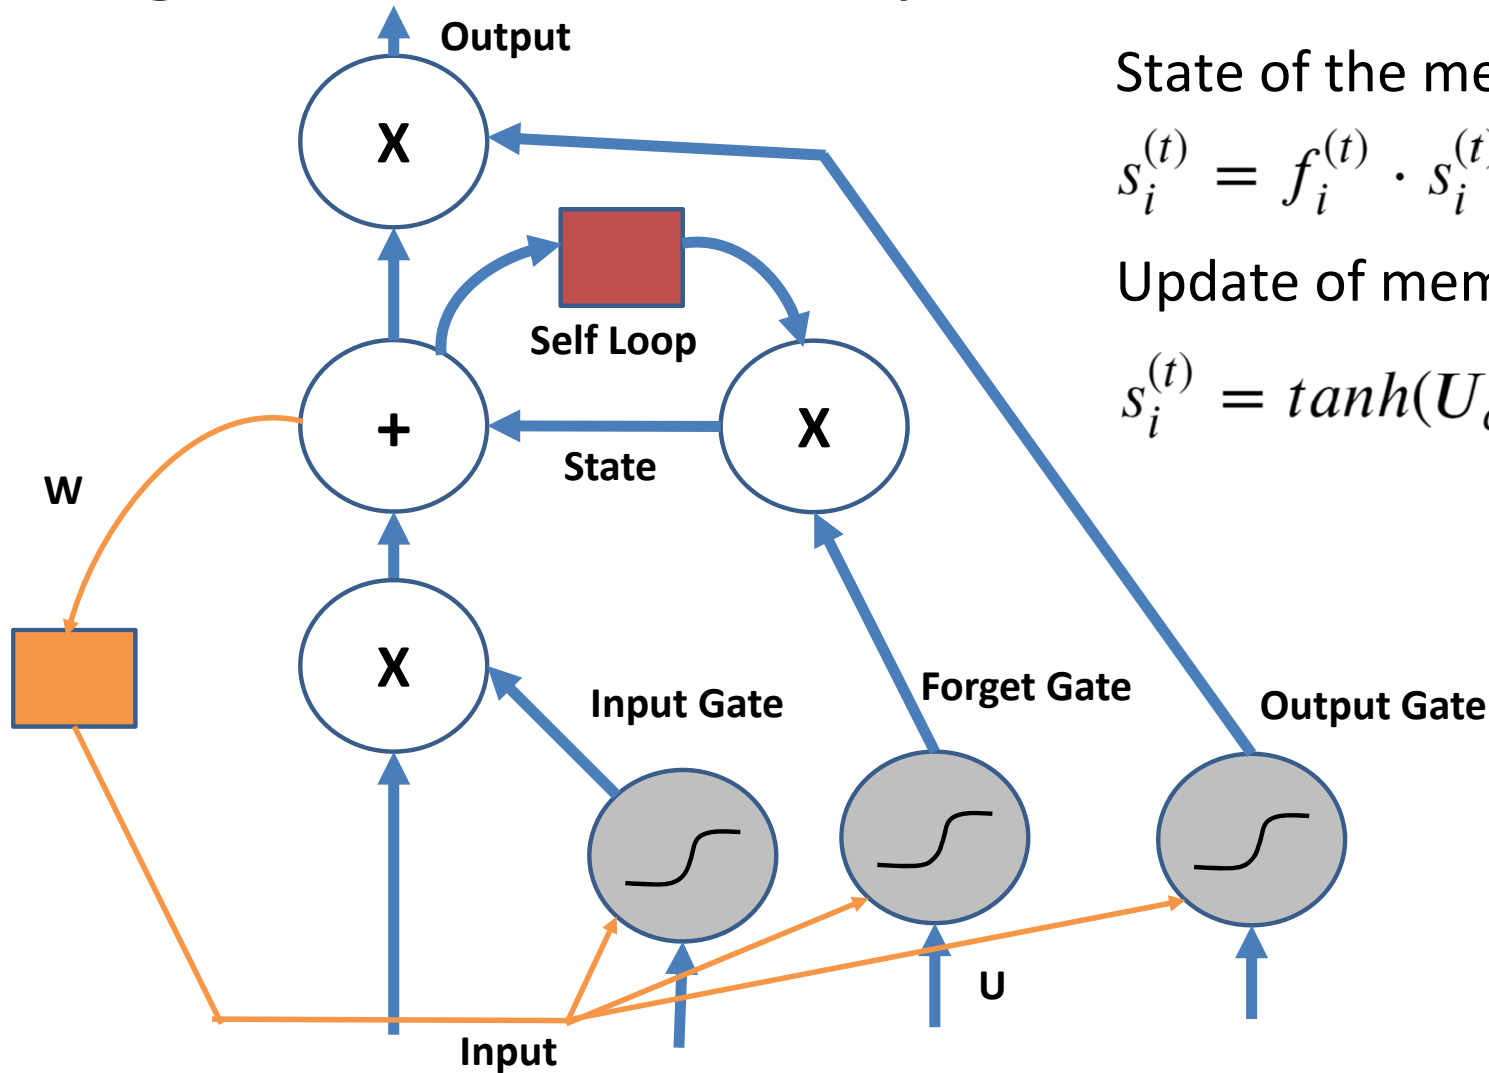
Output of LSTM unit

$$\hat{y}_i^{(t)} = \sigma_i^{(t)} \cdot \tanh(y_i^{(t)})$$

Activation of output gate

$$\sigma_i^{(t)} = \sigma(V_{(o)}s_i^{(t)} + U_{(o)}x_i^{(t)} + W_{(o)}h_i^{(t-1)})$$

# Long-Short Term Memory



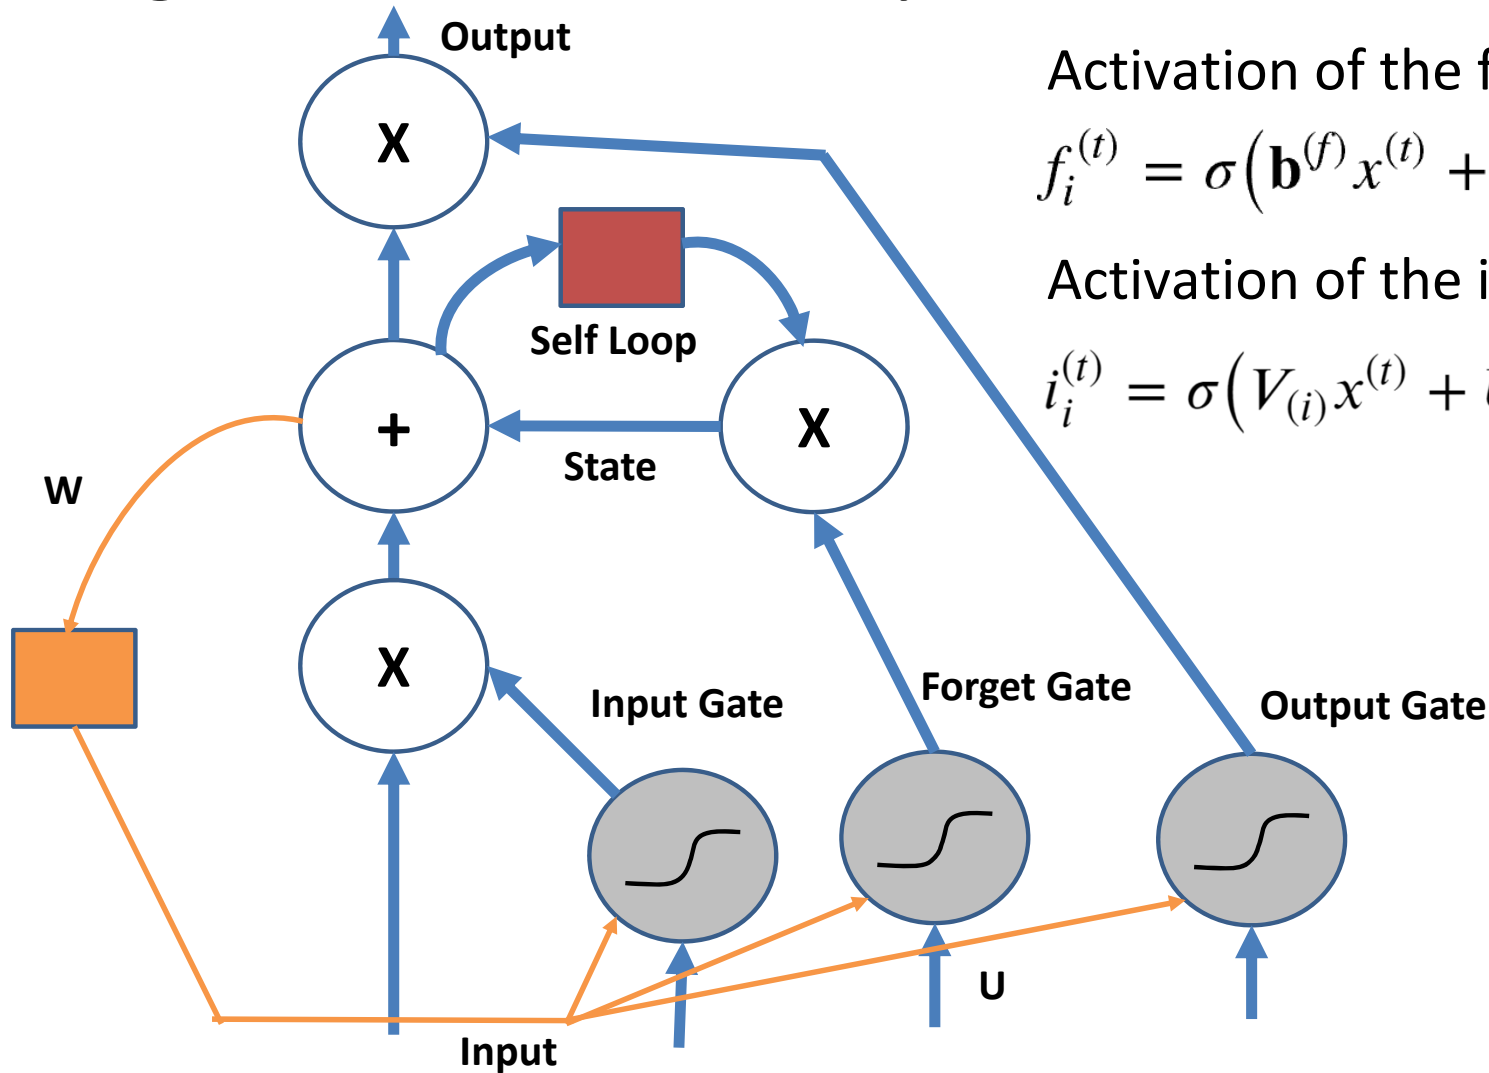
State of the memory loop

$$s_i^{(t)} = f_i^{(t)} \cdot s_i^{(t)} + i_i^{(t)} \cdot s_i^{(t-1)}$$

Update of memory state

$$s_i^{(t)} = \tanh(U_c x_i^{(t)} + W_c h_i^{(t-1)})$$

# Long-Short Term Memory



Activation of the forget gate

$$f_i^{(t)} = \sigma(\mathbf{b}^{(f)} x^{(t)} + \mathbf{U}^{(f)} x^{(t)} + \mathbf{W}^{(f)} h^{(t-1)})$$

Activation of the input gate

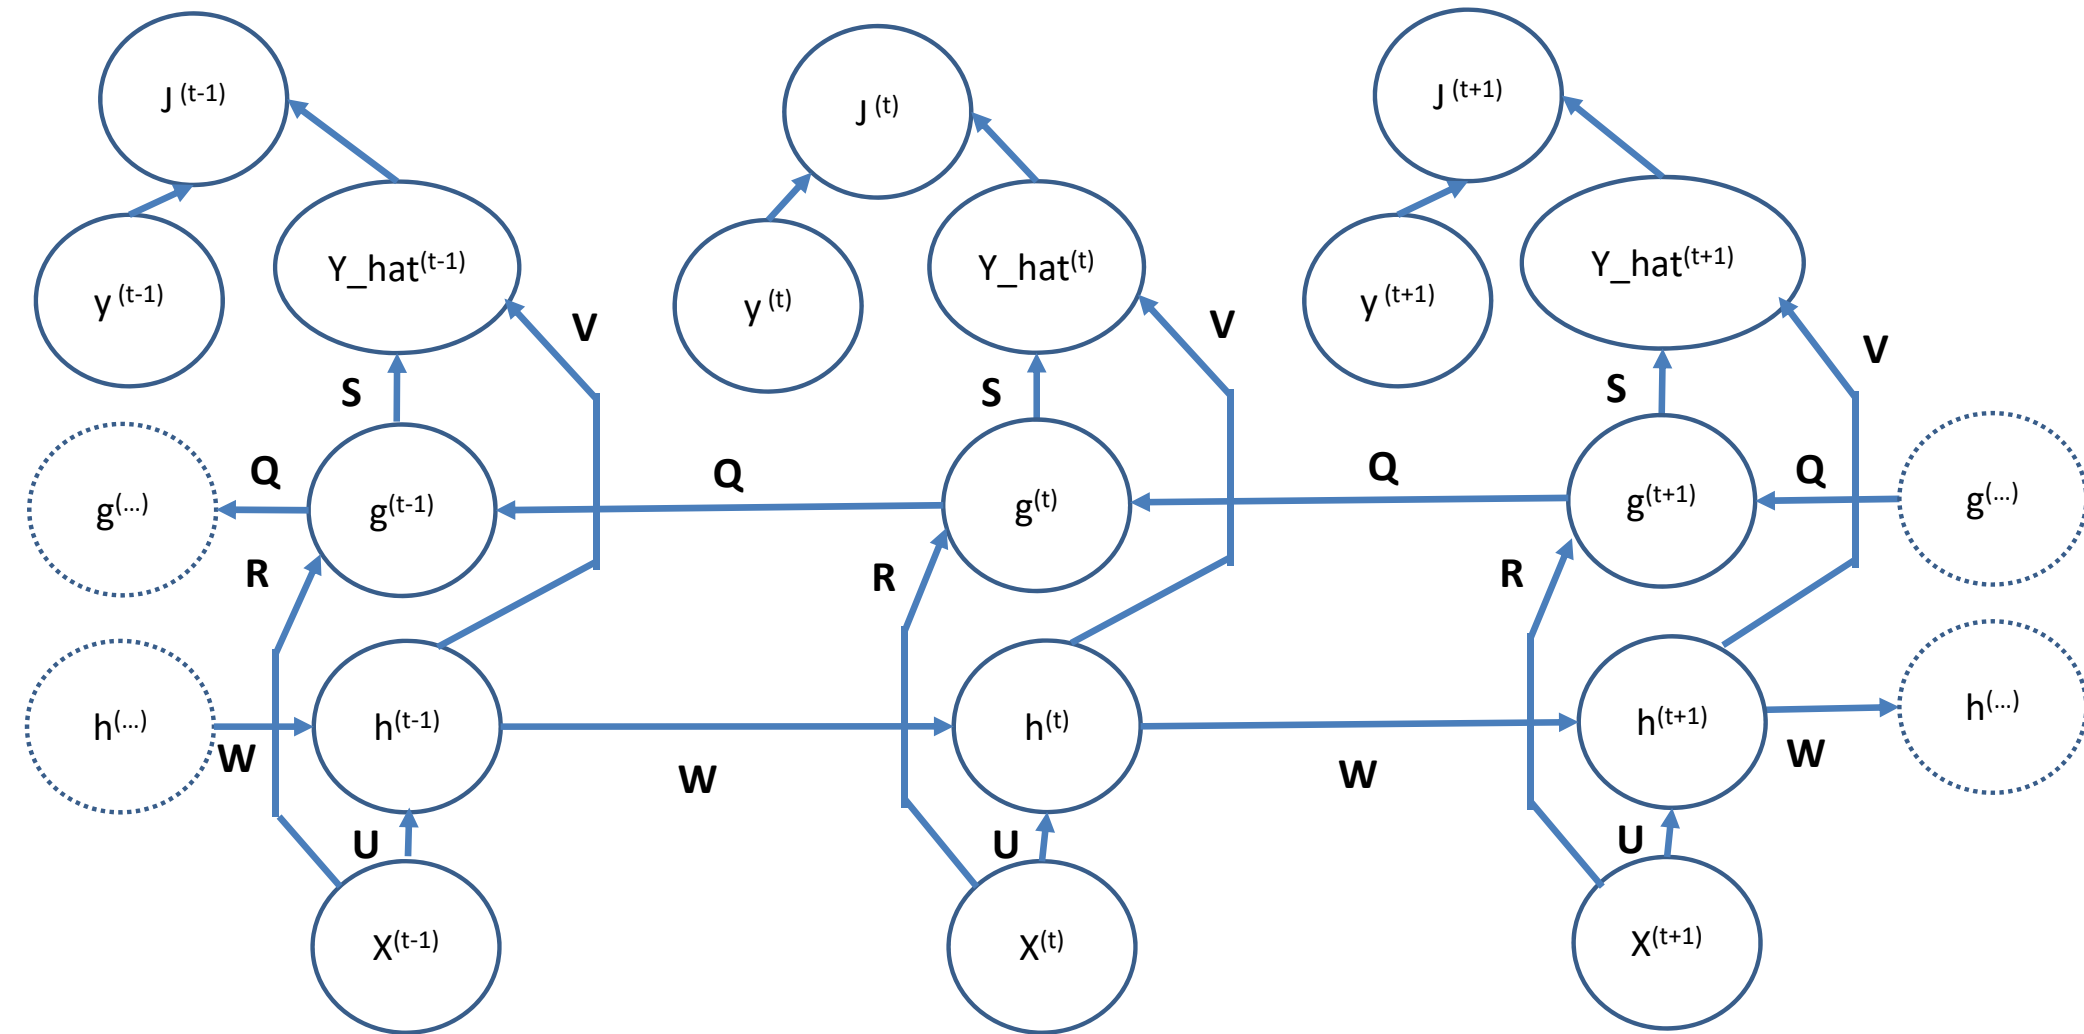
$$i_i^{(t)} = \sigma(V_{(i)} x^{(t)} + U_i x_{(i)}^{(t)} + W_{(i)} h_i^{(t-1)})$$



# Bidirectional RNNs

- Not all sequential relationships are causal
  - A natural language phrase can be parsed in both directions
  - Hand writing recognition can proceed from either end
  - Figure captioning has no preferred direction
- For non-causal sequences we can use **bidirectional RNNs**
- Bidirectional RNNs trained using BPTT in both directions

# Bidirectional RNNs



# Statistical Machine Translation

How to build a statistical model to translate from one language to another?

- Statistical machine translation has a long history, starting in the 1950s
- Goal:
  - Given an **input sequence**,  $\mathbf{X} = x_1, x_2, \dots, x_n$ , of tokens in a first language
  - Find the most probable or **target output sequence**,  $\mathbf{y} = y_1, y_2, \dots, y_m$ , in a second language

# Statistical Machine Translation

How to build a statistical model to translate from one language to another?

- The most probability of the output sequence  $\mathbf{y}$ , given the input sequence  $\mathbf{x}$  follows this relationship

$$p(\mathbf{y}|\mathbf{x}) \propto p(\mathbf{x}|\mathbf{y})p(\mathbf{y})$$

- Where,

$p(\mathbf{x}|\mathbf{y})$  is the **translation model**, or likelihood of sequence  $\mathbf{x}$  given sequence  $\mathbf{y}$

$p(\mathbf{y})$  is the language model, or probability of sequence  $\mathbf{y}$  in the target language

# Statistical Machine Translation

How to build a statistical model to translate from one language to another?

- Find most probably target sequence

$$p(\mathbf{y}|\mathbf{x}) \propto p(\mathbf{x}|\mathbf{y})p(\mathbf{y})$$

- Can find weights for the maximum likelihood expression

$$\mathbf{W} = w_1, w_2, \dots, w_N$$

$$\log p(\mathbf{y}|\mathbf{x}) = \sum_{n=1}^N w_n p(\mathbf{y}|\mathbf{x}) + \log Z(\mathbf{x})$$

- Where  $Z(\mathbf{x})$  is the normalization constant not dependent on  $\mathbf{y}$

# Evaluation of SMT Models

## The BLEU score

- The **Bilingual Translation Understudy** or **BLEU score** is used to compare SMT models.
- BLEU is a comparison between results of an SMT model and results of expert human translators
- The BLEU compares a hypothesis of *i-gram* tuples,  $H(i)$
- The BLEU is the average proportion of *i-gram* matches
- A model with a higher BLEU is considered superior

# Evaluation of SMT Models

The BLEU score

The **BLEU score** for a sequence length N is computed:

$$BLEU = \left\{ \prod_{i=1}^N P(i) \right\}^{\frac{1}{N}}$$

Where,

$$P(i) = \frac{Matched(i)}{H(i)}$$

# Evaluation of SMT Models

## The BLEU score

- $H(i)$  is the number of *i-gram* tuples in each hypothesis  
For hypothesis of length  $n$ , examples are:  $H(1) = n$ ,  $H(2) = n-1$ ,  
 $H(3) = n-2$
- And the number of matches:

$$Matched(i) = \sum_{t_i} \min\{C_h(t_i), \max_j C_{hj}(t_i)\}$$

Where, for *i*-gram tuple  $t_i$

$C_h(t_i)$  = number of times  $t_i$  occurs in the hypothesis to be tested

$C_{hj}(t_i)$  = number of times  $t_i$  occurs in the  $j$ -th reference



# Evaluation of SMT Models

## The BLEU score

- The **BLEU score** is dependent on the sequence length
- A commonly used adjustment is the **brevity penalty,  $\rho$**
- The adjusted BLEU score is then:

$$BLEU_{\rho} = \exp\left(\min\left(0, \frac{n - L}{n}\right)\right) \left\{ \prod_{i=1}^N P(i) \right\}^{\frac{1}{N}}$$

Where:

$n$  = length of the hypothesis

$L$  = length of the reference sequence

# Evaluation of SMT Models

## BLEU score example

- Start with French sentence:  
La voiture est dans l'allée.
- An expert English translation is:  
The car is in the drive.

# Evaluation of SMT Models

## BLEU score example

- Start with the expert English translation:  
The car is in the drive.
- There are  $6 - 1 = 5$ , 2-grams for this sentence:

Reference 2-grams
The car
Car is
Is in
In the
The drive

A perfect MT is identical so:

Matched = 5

$H(2) = 5$

$\rho = 1.0$

BLEU = 100.0%

# Evaluation of SMT Models

## BLEU score example

- A possible MT is:  
The drive contains the car.
- There are  $5 - 1 = 4$ , 2-grams for this sentence:

Reference 2-grams	MT 2-grams	Match
The car	The drive	1
Car is	Drive contains	0
Is in	Contains the	0
In the	The car	1
The drive		

Matched = 2

$H(2) = 5$

$n = 6$

$l = 5$

$\rho = 1.18$

BLEU = 47.3%

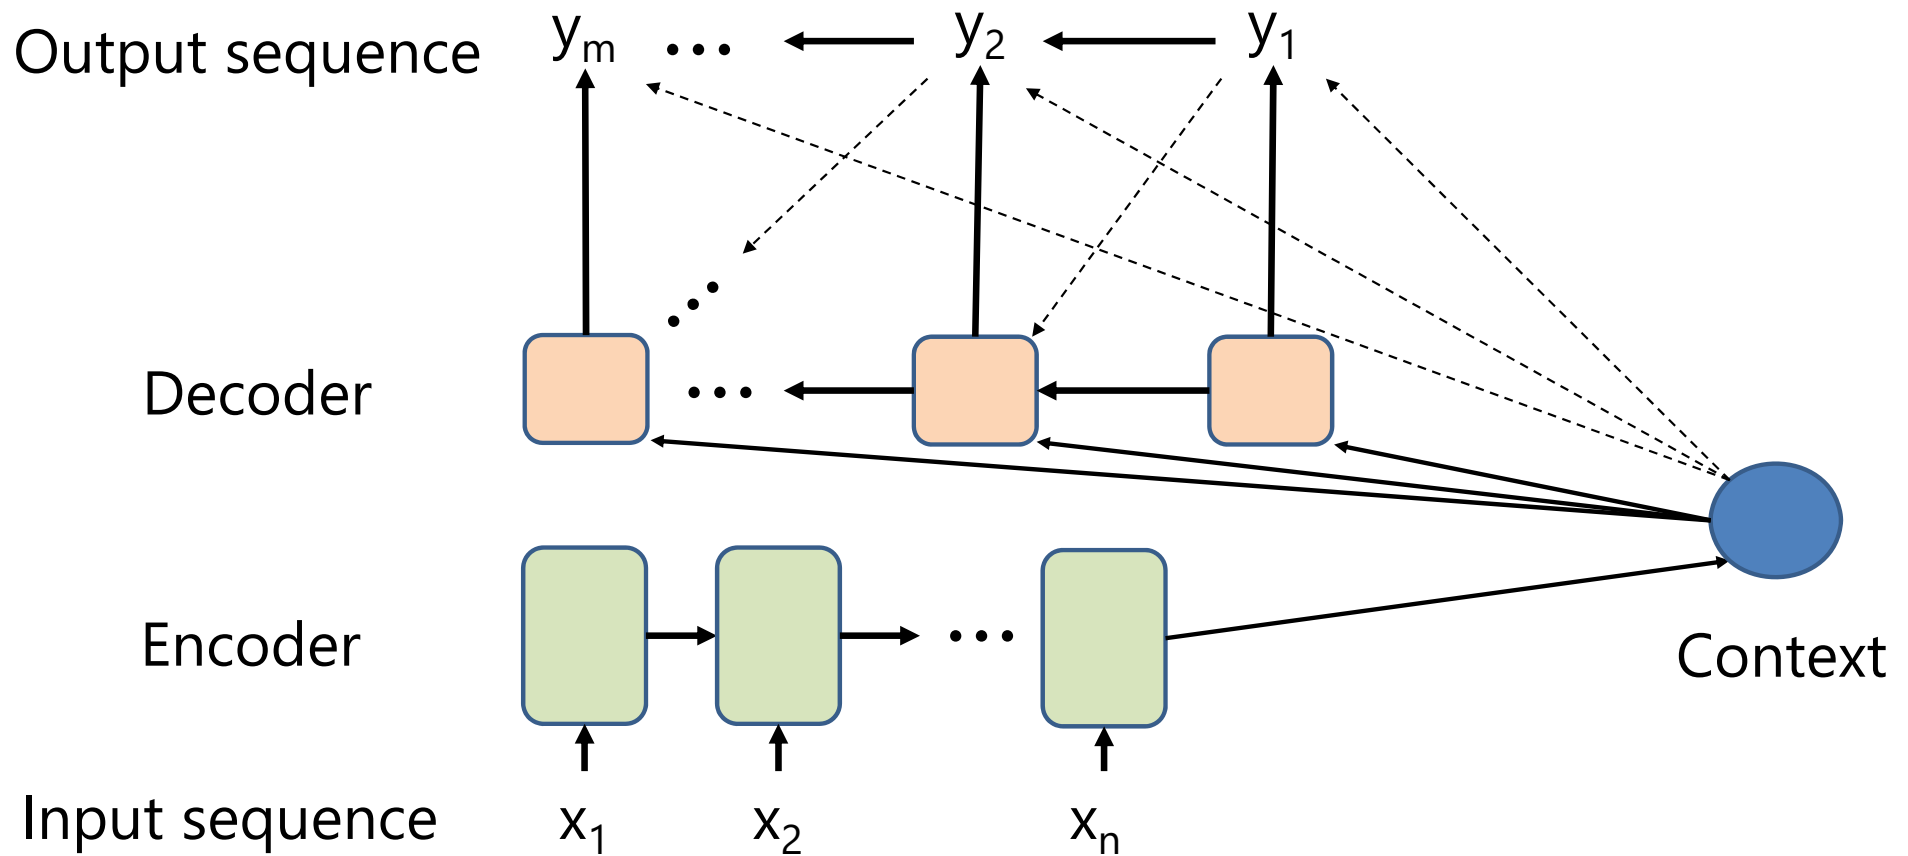
# Neural Machine Translation

Use neural networks as function approximators for machine translators

- Architecture has three components:
  - **Encoder** for the input sequence
  - **Context vector** – a hidden state
  - **Decoder** generates output sequence
- Finds maximum likelihood target sequence

# NMT Architectures

## Encoder-decoder model with context vector



# Neural Machine Translation

Use neural networks as function approximators for machine translators

- The hidden state updates as a function current hidden state and input

$$h_t = f(h_{t-1}, x_t)$$

Where,  $f()$  is the activation function of the LSTM

- The encoder creates a sequence of hidden states given the input sequence

$$h_1, h_2, \dots, h_n = \text{Encoder}_{RNN}(x_1, x_2, \dots, x_n)$$

# Neural Machine Translation

Use neural networks as function approximators for machine translators

- The first layers of a NMT encoder network are a bidirectional RNN (BiRNN)
- The hidden states of the forward layer are:

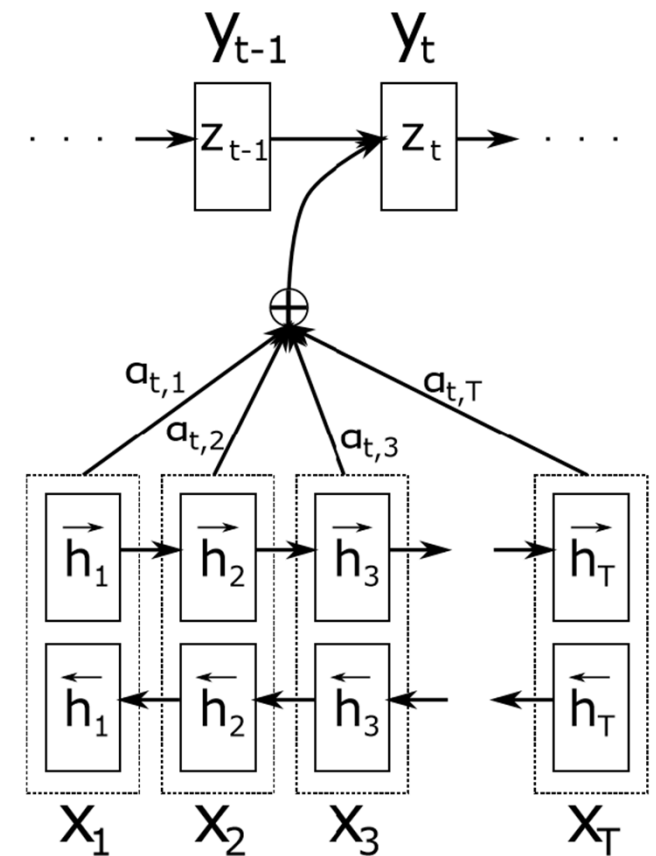
$$\{\vec{h}_1, \vec{h}_2, \dots, \vec{h}_T\}$$

- The hidden states of the reverse layer are:

$$\{\overleftarrow{h}_T, \overleftarrow{h}_{T-1}, \dots, \overleftarrow{h}_1\}$$

- And, the context is:

$$c_t = [\vec{h}_t; \overleftarrow{h}_t]$$





# Neural Machine Translation

Use neural networks as function approximators for machine translators

- Given weights,  $\mathbf{w} = w_1, w_2, \dots, w_N$ , the encoder uses the context and its hidden state

$$p(y_1, y_2, \dots, y_m) = \prod_{t=1}^m p(y_t | \{y, c, h\}_{\leq m})$$

- Taking logs of both sides, gives:

$$\log p(y_1, y_2, \dots, y_m) = \sum_{t=1}^m \log p(y_t | \{y, c, h\}_{\leq m})$$

# Attention Mechanism

- **Context** is the input to decoder
- A **fixed context vector** has limited representation
  - May compress information required for sequence generation
  - Translation accuracy decreases with sequence length
- Need a better representation!
- Use a **context set**:  $\mathcal{C} = \{c_1, c_2, \dots, c_M\}$
- The **weights** of the context give **attention** to the correct sequence in the decoder

# Attention Mechanism

How to compute the **attention weights**?

- Compute **attention score** for context  $c_i$  as a function of pervious hidden state, context and previous attention weights:

$$e_i^t = f_{ATT}(h_{t-1}, c_i, \{\alpha_j^{t-1}\}_{j=1}^M)$$

- The attention weights are updated:


$$\alpha_i^t = \frac{\exp(e_i^t)}{\sum_{j=1}^M \exp(e_j^t)}$$

- The attention weights are the **probability that decoder should attend to the context  $c_i$**

# Attention Mechanism

## Example of attention weights

- Attention weights can be readily visualized
- For the input sequence
- And, the output sequence
- **Probabilities of attending the output sequence given the input sequence**



# Attention Mechanism

How to compute the **attention weights**?

- The scores used to update the context weight are a function of the current context set and the attention weights:

$$e_i^t = \phi(\{c_i^t\}_{i=1}^M, \{\alpha_i^t\}_{i=1}^M)$$

- One possibility is to use linear combination of weighted context:

$$e_i^t = \phi(\{c_i^t\}_{i=1}^M, \{\alpha_i^t\}_{i=1}^M) = \sum_{i=1}^M \alpha_i c_i$$

# Attention Mechanism

How to compute the **attention weights**?

- Use a neural network as a function approximator
- The neural network has parameters,  $\Theta$ , which maximize the log-likelihood

$$\mathcal{L}(D, \Theta) = \frac{1}{N} \sum_{n=1}^N \log p(y^n | x^n, \Theta)$$

- Where the data are the N input and output sequences of the training data;  $D = \{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$
- For linear model, can compute all the derivatives and use backpropagation to train neural network with stochastic gradient decent

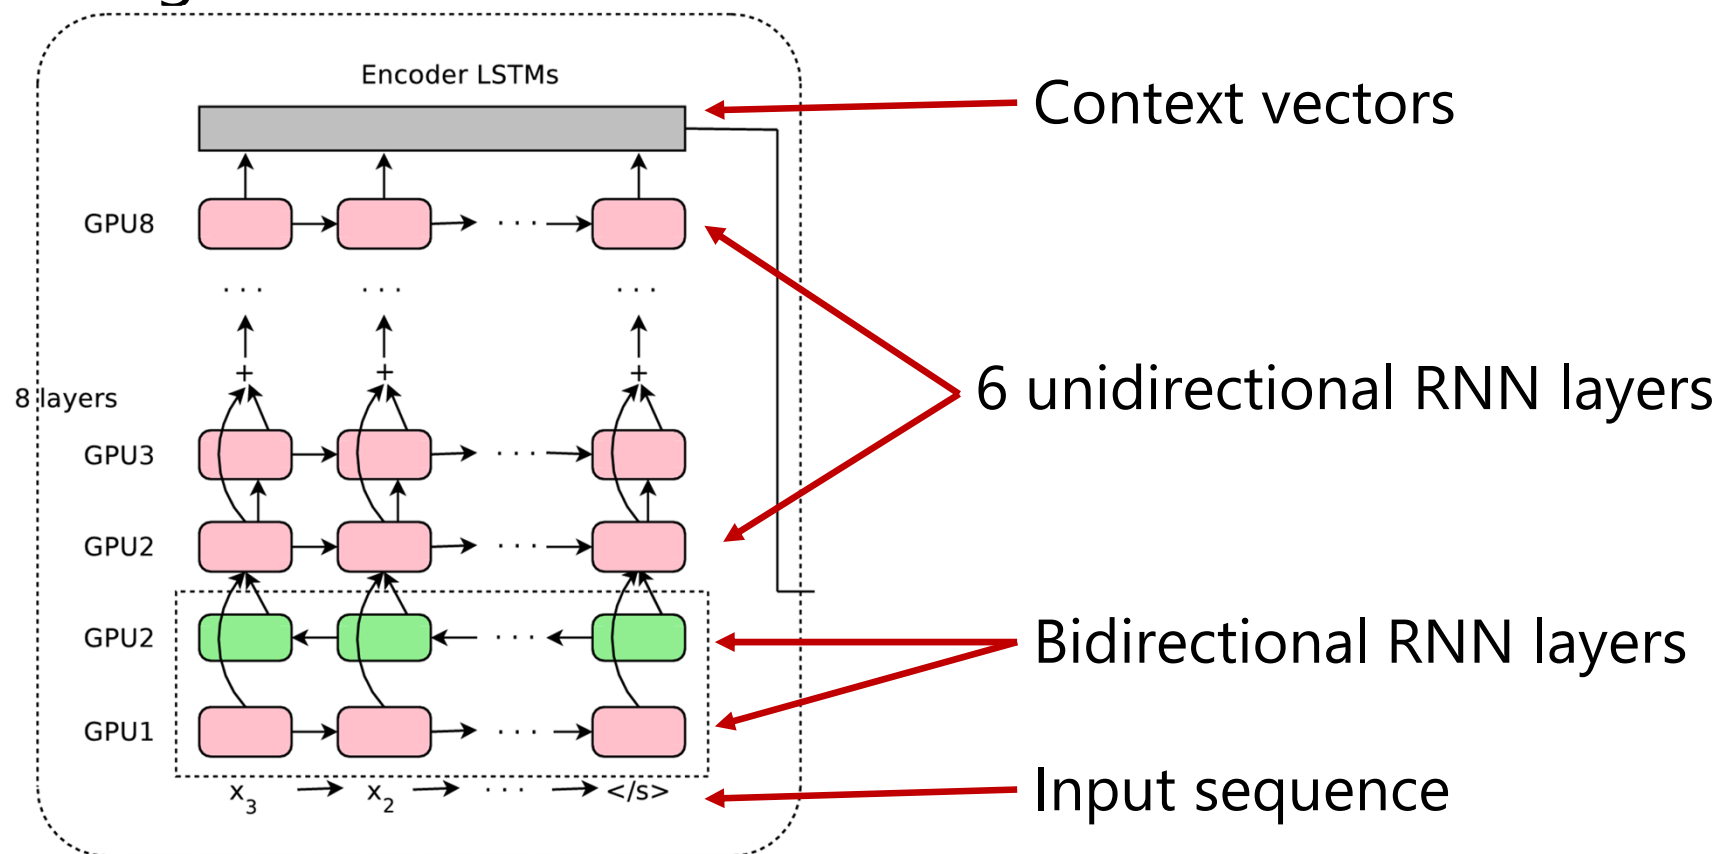
# Beam Search in Decoder

Use the beam search method to find the most probable output sequence

- Given the attention probabilities,  $\alpha_i^t$ , how can the decoder find the most probable output sequence?
- Use a search heuristic known as beam search
- Beam search is a classic AI method originating in the 1970s
- Beam search is a **breath first search method**
  - Search along k highest probability paths
  - Uses less memory than full breath search
- Beam search used to find the sequence with the highest probability or likelihood

# Google Translate

## Google Translate architecture: Encoder

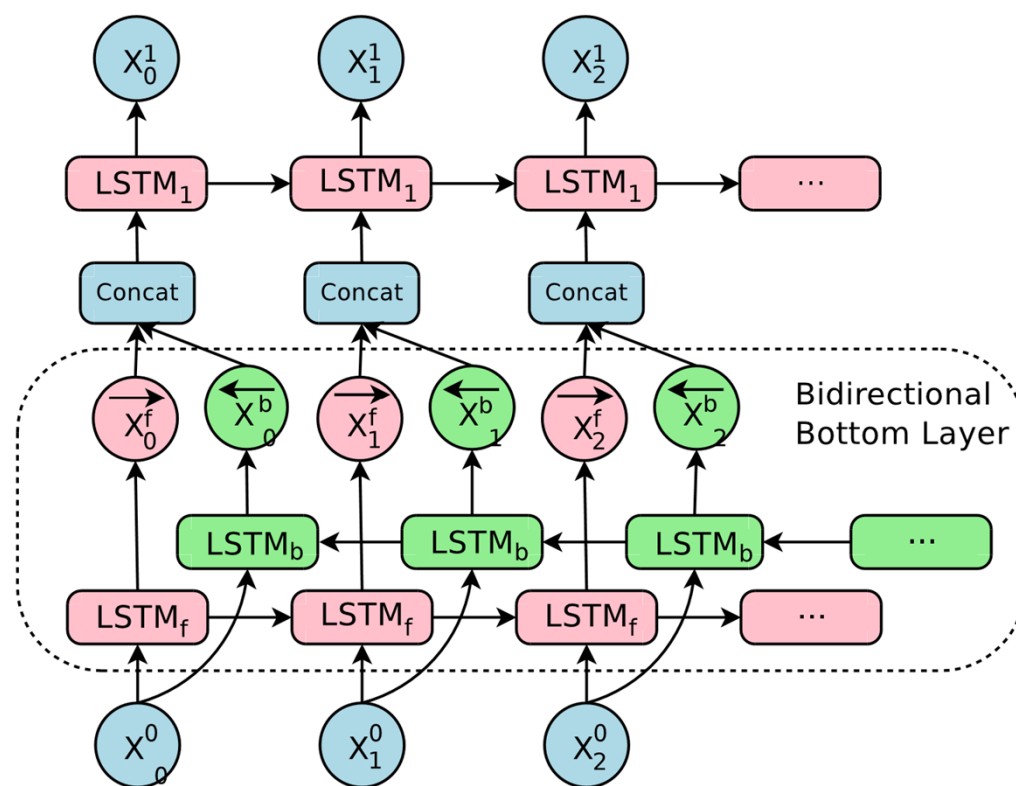




# Google Translate

## Google Translate architecture

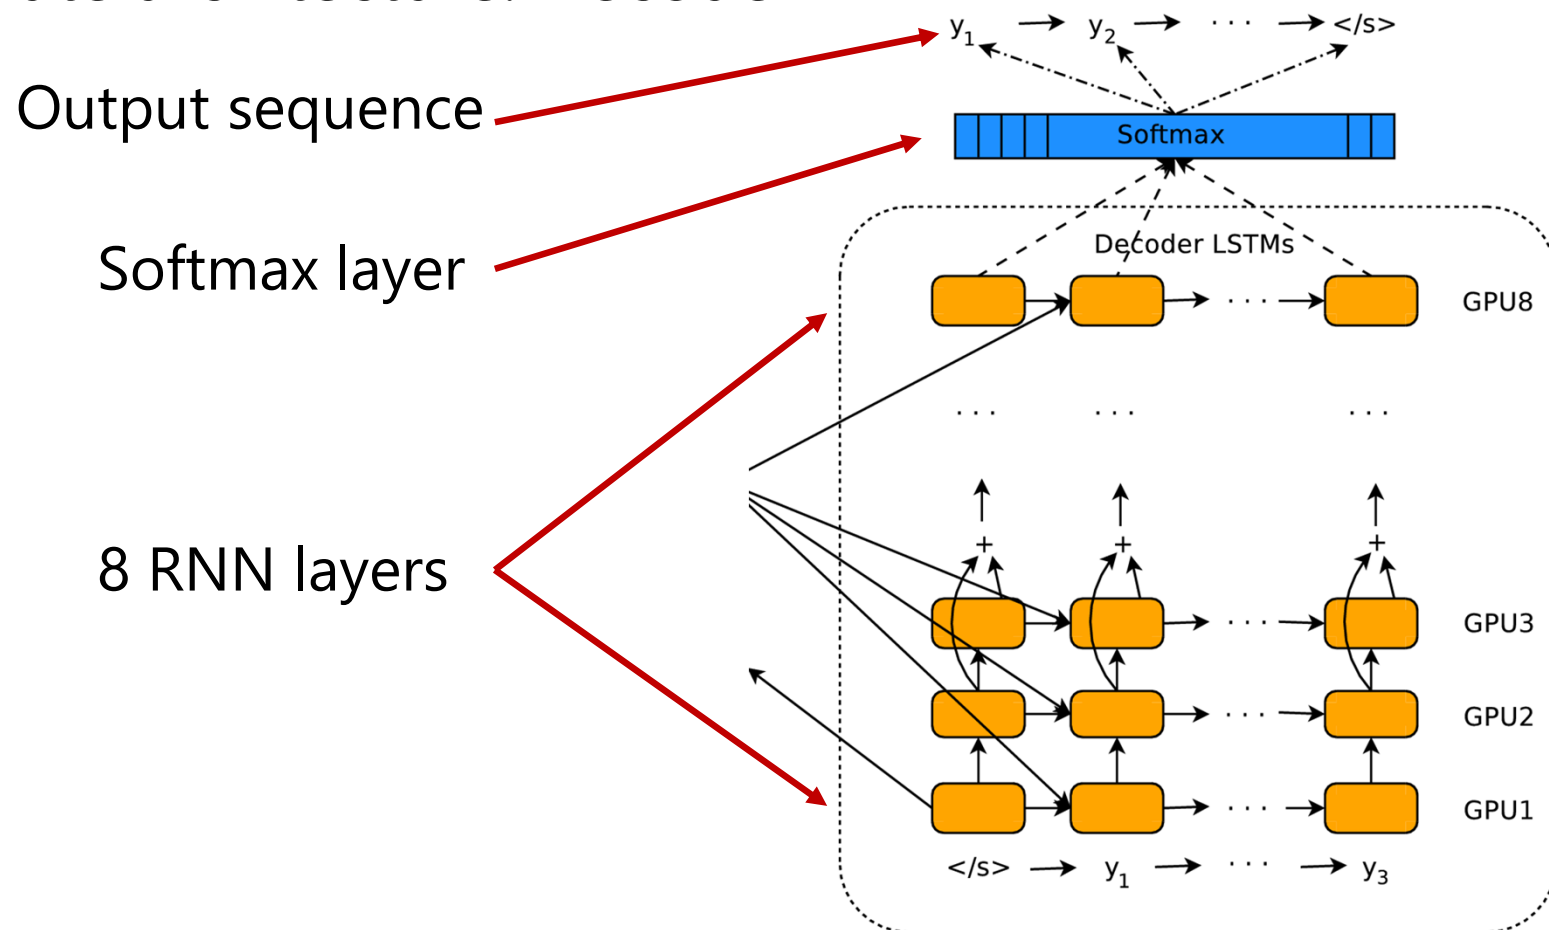
### Details of the biRNN layer



[Wu, et. al, Google's Neural Machine Translation System: Bridging the Gap between Human and Machine Translation, 2016](#)

# Google Translate

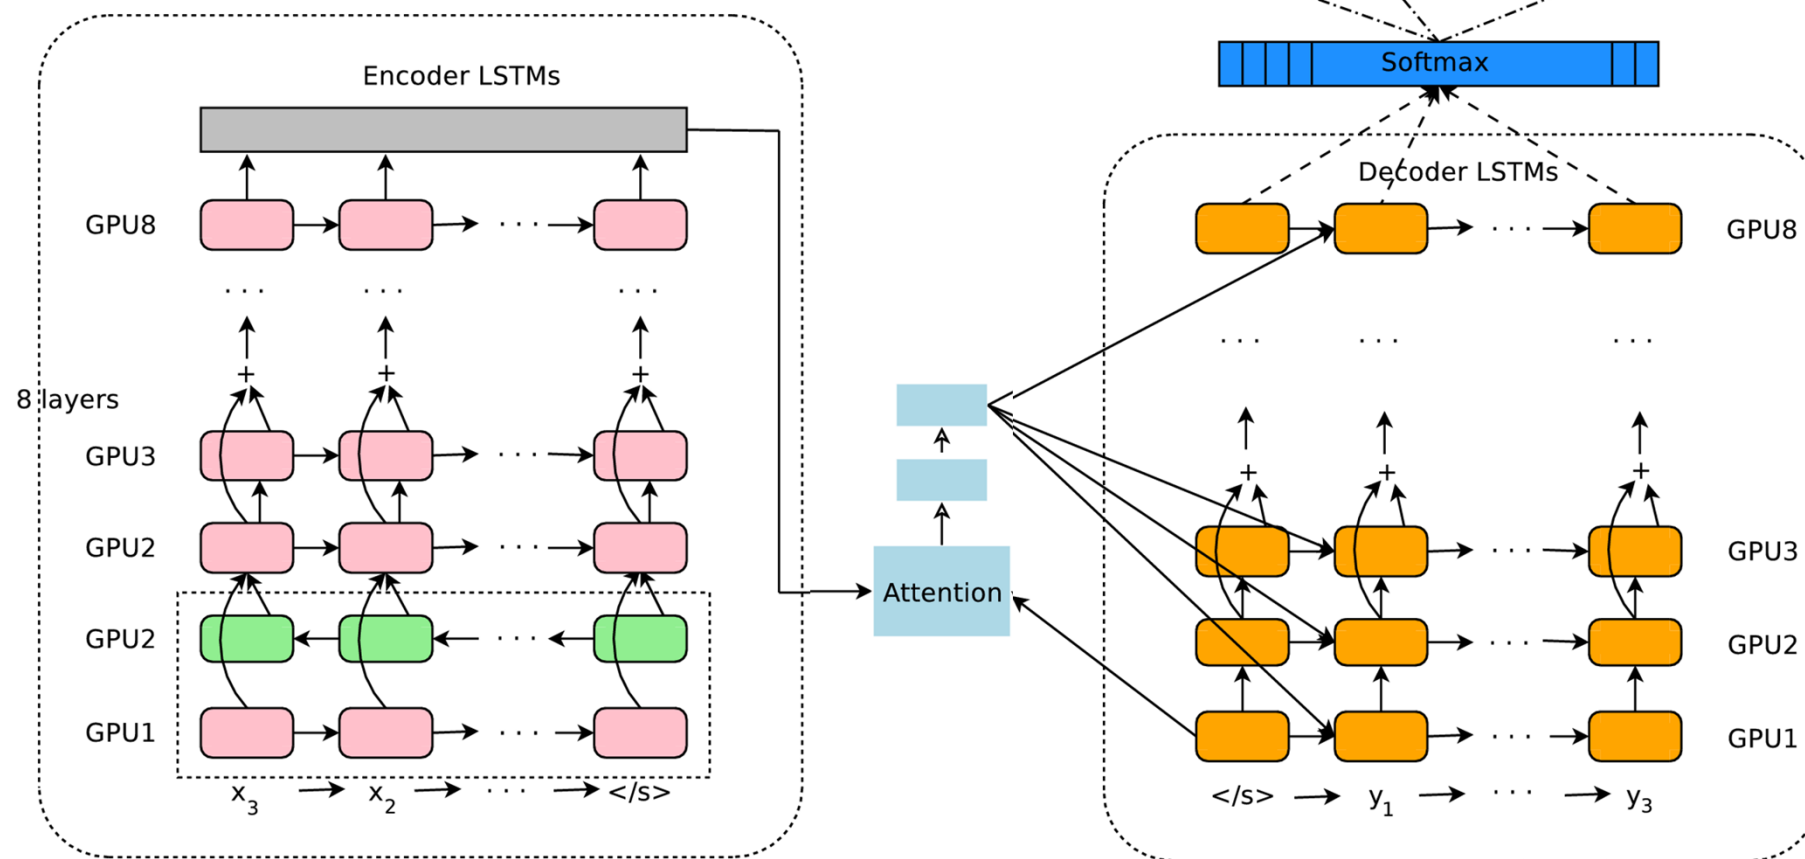
## Google Translate architecture: Decoder



[Wu, et. al, Google's Neural Machine Translation System: Bridging the Gap between Human and Machine Translation, 2016](#)

# Google Translate

## Google Translate architecture



[Wu, et. al, Google's Neural Machine Translation System: Bridging the Gap between Human and Machine Translation, 2016](#)

# Google Translate

## Beam search for the attention mechanism

- Need a decoder which does not favor particular output sequence
- Standard likelihood score needs adjustments
  - Do not want to favor short sequences resulting from the summation of negative log likelihoods
  - Need to ensure that the entire input sequence is reflected in the output sequence

# Google Translate

## Beam search for the attention mechanism

The modified likelihood score becomes:

$$\mathcal{L}(D, \Theta) = \frac{1}{lp(Y)} \frac{1}{N} \sum_{n=1}^N \log p(y^n | x^n, \Theta) + cp(D)$$

$$lp(Y) = \frac{(5 + |Y|)^\gamma}{(5 + 1)^\gamma}$$

$$cp(D) = \beta \sum_{i=1}^{|X|} \log \left( \min \left( \sum_{j=1}^{|Y|} \alpha_{i,j}, 1.0 \right) \right)$$

Where


$\gamma$  = strength of the normalization

$\beta$  = extent to which translations that fully cover the source sentence are favored

# Google Translate

## How to train deep LSTM RNNs?

- Stacking many LSTM layers is problematic
- Gradient vanishes
- Solution: use residual connections!



# Google Translate

## How to train deep LSTM RNNs?

- The following relations govern the update between the i-th and i+1-th standard LSTM layers:

$$h_t^i, m_t^i = LSTM_i(h_{t-1}^i, m_{t-1}^i, x_t^{i-1}, \mathbf{W}^i)$$

$$x_t^i = h_t^i$$

$$h_t^{i+1}, m_t^{i+1} = LSTM_{i+1}(h_{t-1}^{i+1}, m_{t-1}^{i+1}, x_t^i, \mathbf{W}^{i+1})$$

Where, for the i-th layer:

$m_i$  = the LSTM memory

$h_i$  = hidden state

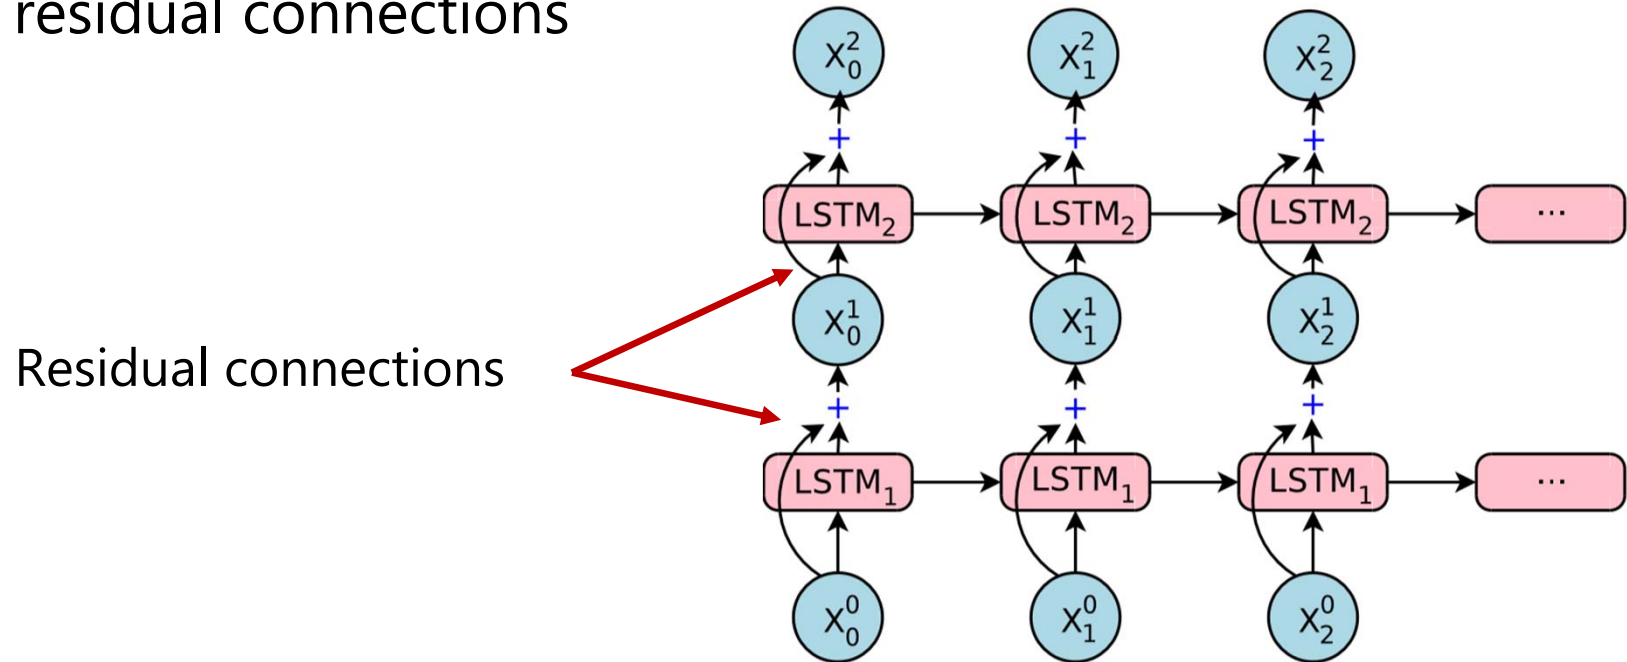
$x_i$  = input

$\mathbf{W}_i$  = weight tensor

# Google Translate

## Google Translate architecture: residual connections

- How to deal with vanishing gradient in deep RNN?
- Use residual connections





# Google Translate

## How to train deep LSTM RNNs?

- The following relations govern the update between the  $i$ -th and  $i+1$ -th standard LSTM layers:

$$h_t^i, m_t^i = LSTM_i(h_{t-1}^i, m_{t-1}^i, x_t^{i-1}, \mathbf{W}^i)$$

$$x_t^i = h_t^i + x_t^{i-1} \quad \text{< residual connection}$$

$$h_t^{i+1}, m_t^{i+1} = LSTM_{i+1}(h_{t-1}^{i+1}, m_{t-1}^{i+1}, x_t^i, \mathbf{W}^{i+1})$$

# Google Translate

How well does it work?

Compare BLEU scores with and without reinforcement learning fine tuning

Table 6: Single model test BLEU scores, averaged over 8 runs, on WMT En→Fr and En→De

Dataset	Trained with log-likelihood	Refined with RL
En→Fr	38.95	39.92
En→De	24.67	24.60

Simple NMT models have BLEU scores around 30 for En-> FR

[Wu, et. al, Google's Neural Machine Translation System: Bridging the Gap between Human and Machine Translation, 2016](#)

# Google Translate

How well does it work?

Using an ensemble of models helps:

Table 7: Model ensemble results on WMT En→Fr (newstest2014)

Model	BLEU
WPM-32K (8 models)	40.35
RL-refined WPM-32K (8 models)	41.16
LSTM (6 layers) [31]	35.6
LSTM (6 layers + PosUnk) [31]	37.5
Deep-Att + PosUnk (8 models) [45]	40.4